
C335

Computer Structures

Memory Hierarchies (III)

Dr. Liqiang Zhang

Department of Computer and Information Sciences


Adapted from Morgan Kaufmann and others

Block Size Tradeoff (1/3)



❑ Benefits of Larger Block Size

- Spatial Locality: if we access a given word, we're likely to access other nearby words soon
- Very applicable with Stored-Program Concept: if we execute a given instruction, it's likely that we'll execute the next few as well
- Works nicely in sequential array accesses too



❑ Drawbacks of Larger Block Size

- Larger block size means **larger miss penalty**
 - on a miss, takes longer time to load a new block from next level
- If block size is too big relative to cache size, then there are too few blocks
 - Result: miss rate goes up

❑ In general, we want to minimize

Average Memory Access Time (AMAT)

$$= \text{Hit Time} \times \text{Hit Rate} + \text{Miss Penalty} \times \text{Miss Rate}$$

Block Size Tradeoff (3/3)



- ❑ Hit Rate = % of requests that are found in current level cache
- ❑ Miss Rate = $1 - \text{Hit Rate}$
- ❑ Hit Time = time to find and retrieve data from current level cache
- ❑ Miss Penalty = average time to retrieve data on a current level miss (includes the possibility of misses on successive levels of memory hierarchy)




Extreme Example: One Big Block



Valid Bit	Tag	Cache Data
<input type="checkbox"/>	<input type="text"/>	<input type="text" value="B 3"/> <input type="text" value="B 2"/> <input type="text" value="B 1"/> <input type="text" value="B 0"/>

- ❑ Cache Size = 4 bytes Block Size = 4 bytes
 - Only ONE entry (row) in the cache!
- ❑ If item accessed, likely accessed again soon
 - But unlikely will be accessed again immediately!
- ❑ The next access will likely to be a miss again
 - Continually loading data into the cache but discard data (force out) before use it again
 - Nightmare for cache designer: **Ping Pong Effect**

Block Size Tradeoff Conclusions



A Summary on Sources of Cache Misses



❑ Compulsory (cold start, first reference): first access to a block

- “Cold” fact of life: not a whole lot you can do about it, although larger block sizes can reduce the number of misses.

❑ Conflict (collision):

- Multiple memory locations mapped to the same cache location
 - Solution 1: increase cache size (index longer)
 - Solution 2: increase associativity

❑ Capacity:

- Cache cannot contain all blocks accessed by the program
 - Solution: increase cache size

Fully Associative Cache (1/3)



❑ Memory address fields:

- Tag: same as before
- Offset: same as before
- Index: nonexistent!

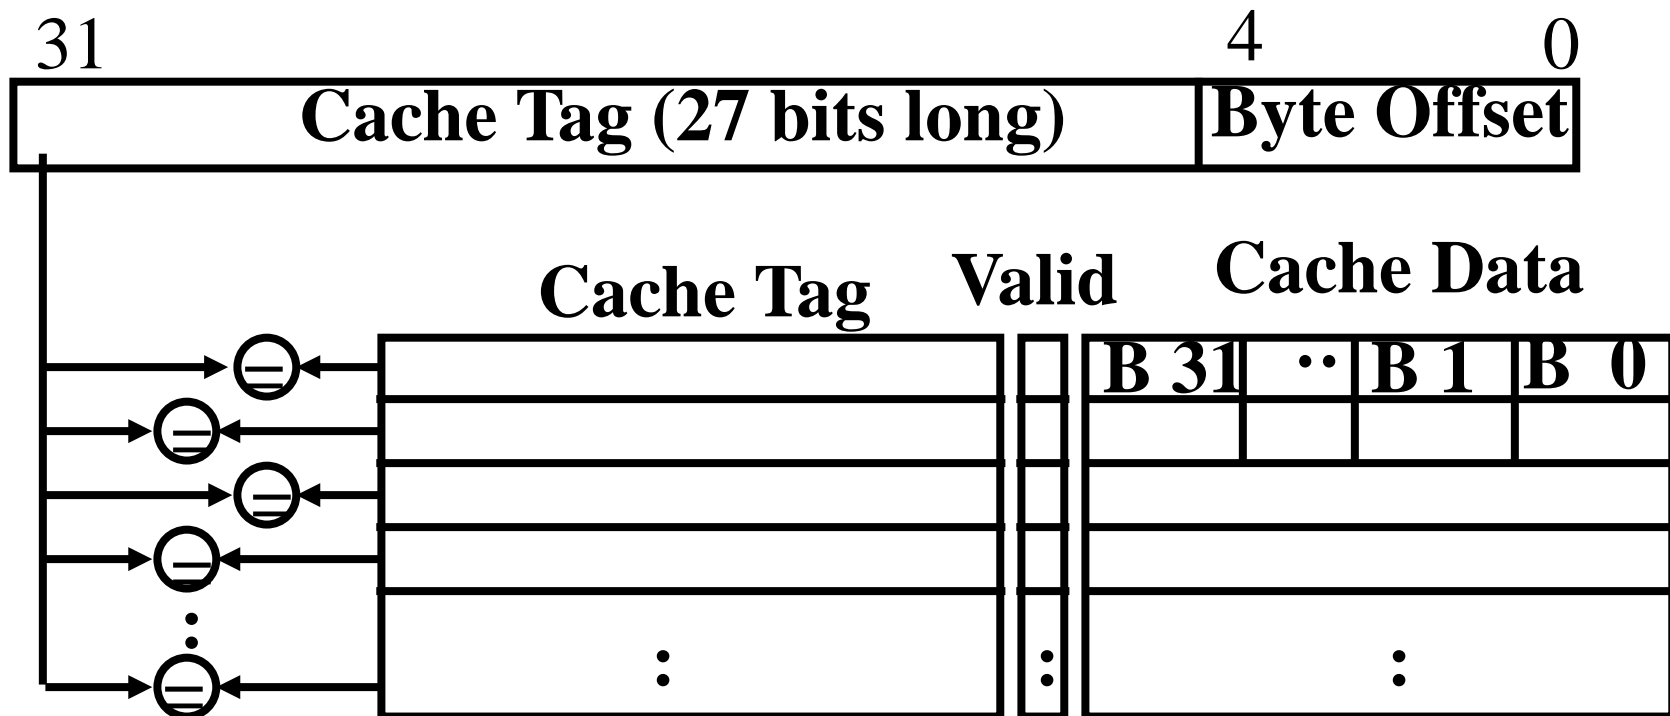
❑ What does this mean?

- no “rows”: any block can go anywhere in the cache
- must compare with all tags in entire cache to see if data is there

Fully Associative Cache (2/3)



- ❑ Fully Associative Cache (e.g., 32 B block)
 - compare tags in parallel



The whole cache could be taken as one single row (the index only has one value – 0)

Fully Associative Cache (3/3)



❑ Benefit of Fully Assoc Cache

- No Conflict Misses (since data can go anywhere)

❑ Drawbacks of Fully Assoc Cache

- Need hardware comparator for every single entry: if we have a 64KB of data in cache with 4B entries, we need 16K comparators: infeasible

N-Way Set Associative Cache (1/3)



❑ Memory address fields:

- **Tag**: same as before
- **Offset**: same as before
- **Index**: points us to the correct “set”

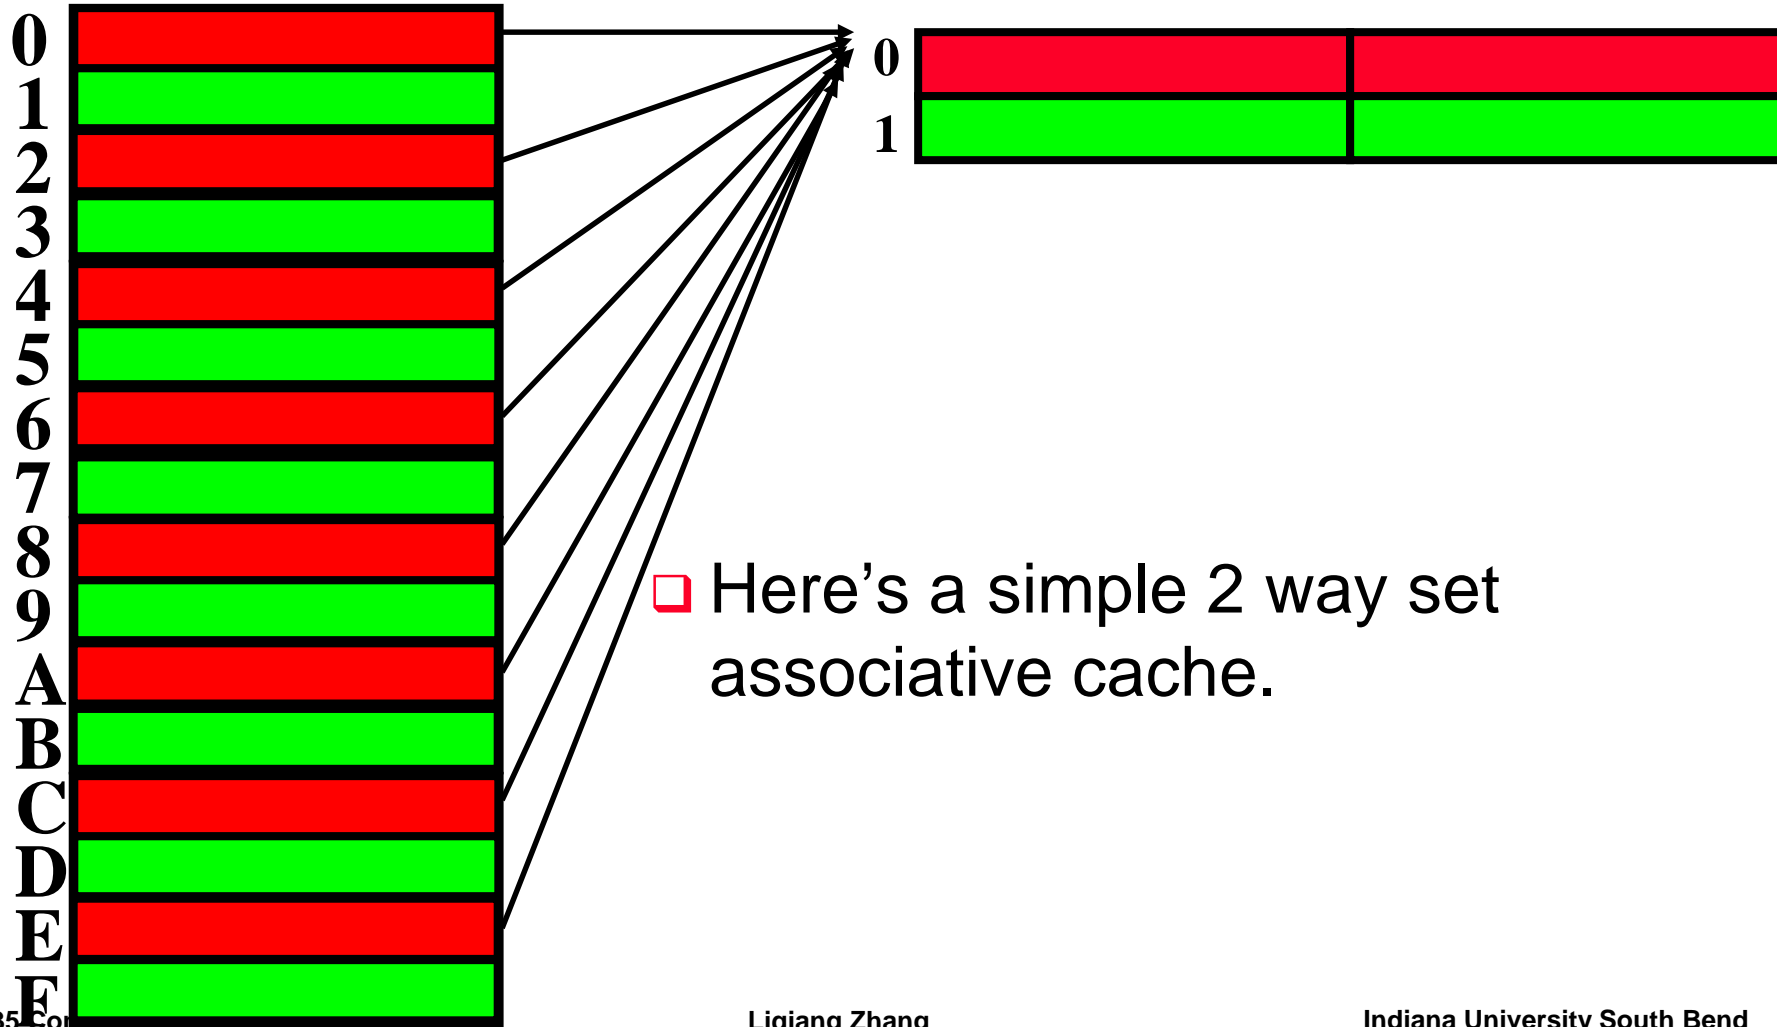
❑ So what's the difference?

- each set contains multiple blocks
- once we've found correct set, must compare with all tags in that set to find our data

Associative Cache Example

Memory

Address Memory



N-Way Set Associative Cache (2/3)



❑ Basic Idea

- cache is direct-mapped w/respect to sets
- each set is fully associative
- basically N direct-mapped caches working in parallel: each has its own valid bit and data

❑ Given memory address:

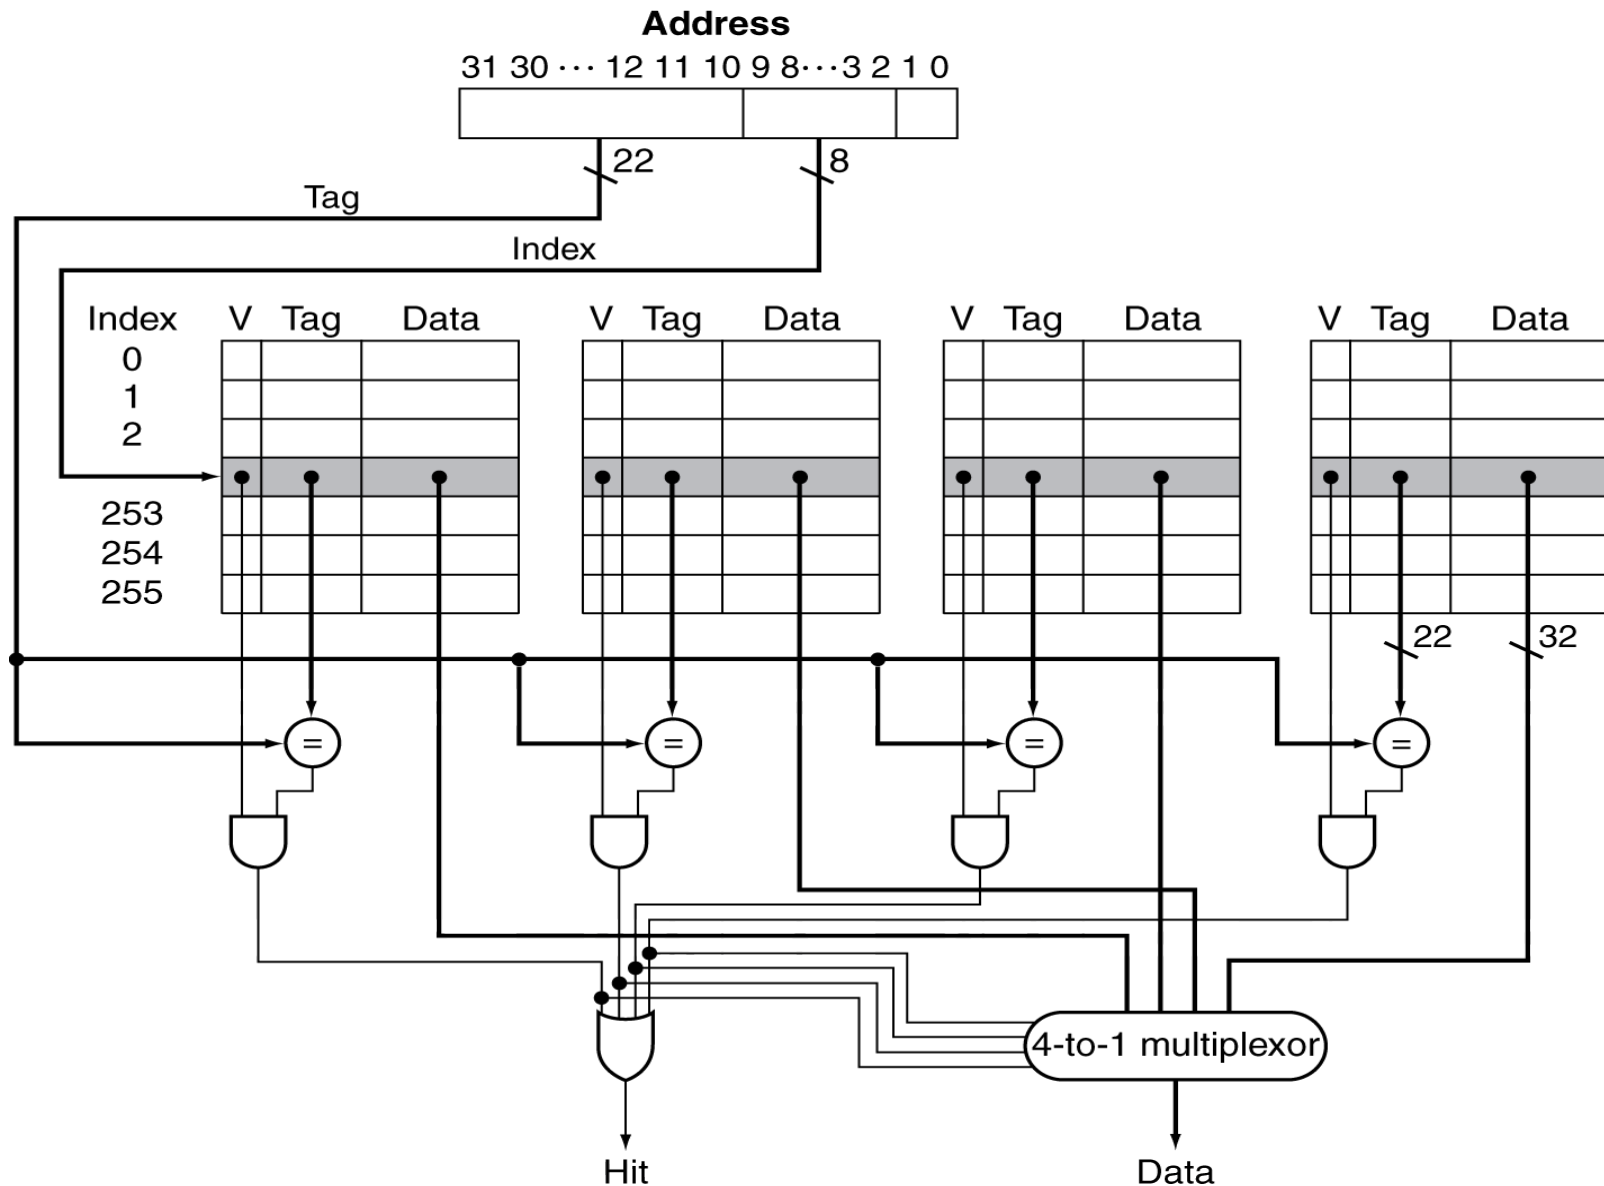
- Find correct set using Index value.
- Compare Tag with all Tag values in the determined set.
- If a match occurs, hit!, otherwise a miss.
- Finally, use the offset field as usual to find the desired data within the block.

N-Way Set Associative Cache (3/3)



- ❑ What's so great about this?
 - even a 2-way set assoc cache avoids a lot of conflict misses
 - hardware cost isn't that bad: only need N comparators
- ❑ In fact, for a cache with M blocks,
 - it's Direct-Mapped if it's 1-way set assoc
 - it's Fully Assoc if it's M -way set assoc
 - so these two are just special cases of the more general set associative design

4-Way Set Associative Cache Circuit



Example: fully associative cache (1/2)



Here is a series of address references given as **word addresses**: 2, 3, 6, 10, 7, 12, 2, 18, 11, and 3. Assuming a **fully associative cache** with 8 **two-word blocks** (total size 16 words) that is initially empty, label each reference in the list as a hit or a miss and show the final contents of the cache. (**note: the basic unit is word!**)

Tag	Data (two-word block)	Tag	Data (two-word block)	Tag	Data (two-word block)	Tag	Data (two-word block)

in one row

Tag	Data (two-word block)	Tag	Data (two-word block)	Tag	Data (two-word block)	Tag	Data (two-word block)

2 (00000010), 3(00000011), 6(00000110), 10(00001010), 7(00000111),
12(00001100), 2(00000010), 18(00010010), 11(00001011), 3(00000011)

Example: fully associative cache (2/2)



Tag	Data (two-word block)	Tag	Data (two-word block)	Tag	Data (two-word block)	Tag	Data (two-word block)
0000001	[2, 3]	0000011	[6, 7]	0000101	[10, 11]	00001110	[12, 13]

in one row

Tag	Data (two-word block)	Tag	Data (two-word block)	Tag	Data (two-word block)	Tag	Data (two-word block)
0001001	[18, 19]						

2 (00000010) → miss, 3(00000011) → hit, 6(00000110) → miss,
10(00001010) → miss, 7(00000111) → hit, 12(00001100) → miss,
2(00000010) → hit, 18(00010010) → miss, 11(00001011) → hit,
3(00000011) → ?

Example: 2-way set associative cache (1/2)

Here is a series of address references given as **word addresses**: 2, 3, 6, 10, 7, 12, 2, 18, 11, and 3. Assuming a **2-way set associative cache** with 8 **two-word blocks** (total size 16 words) that is initially empty, label each reference in the list as a hit or a miss and show the final contents of the cache. (**note: the basic unit is word!**)

Index	Tag	Data (two-word block)	Tag	Data (two-word block)
00				
01				
10				
11				

2 (00000**010**), 3(00000**011**), 6(00000**110**), 10(0000**1010**), 7(00000**111**),
12(0000**1100**), 2(00000**010**), 18(000**10010**), 11(0000**1011**), 3(00000**011**)

Example: 2-way set associative cache (2/2)



Here is a series of address references given as word addresses: 2, 3, 6, 10, 7, 12, 2, 18, 11, and 3. Assuming a 2-way set associative cache with 8 two-word blocks (total size 16 words) that is initially empty, label each reference in the list as a hit or a miss and show the final contents of the cache. (note: the basic unit is word!)

Index	Tag	Data (two-word block)	Tag	Data (two-word block)
00				
01	00000	[2, 3]	00001	[10, 11]
10	00001	[12, 13]		
11	00000	[6, 7]		

2 (00000010) → miss, 3(00000011) → hit, 6(00000110) → miss,
10(00001010) → miss, 7(00000111) → hit, 12(00001100) → miss,
2(00000010) → hit, 18(00010010) → ?
11(00001011) → ?
3(00000011) → ?

Cache Shapes (for 16 blocks)



Direct-mapped

2-way set-associative

4-way set-associative

8-way set-associative

--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--

Fully associative

Cache Block Replacement: Need to Make a Decision!



❑ Direct Mapped Cache:

- Each memory location can only be mapped to 1 cache location
- No need to make any decision :-)
 - Current item replaces the previous item in that cache location

❑ N-way Set Associative Cache:

- Each memory location have a choice of N cache locations

❑ Fully Associative Cache:

- Each memory location can be placed in ANY cache location

❑ Cache miss in an N-way Set Associative or Fully Associative Cache:

- Bring in new block from memory
- Remove a cache block to make room for the new block
- =====> We need to make a decision on which block to replace!

Cache Block Replacement Policy




- ❑ Random Replacement:
 - Hardware randomly selects a cache item and replaces it
- ❑ Least Recently Used (LRU):
 - Hardware keeps track of the access history
 - Replace the entry that has been unused for the longest time
- ❑ Others?

Cache Write Policy: Write Through versus Write Back



- ❑ Cache read is much easier to handle than cache write:
 - Instruction cache is much easier to design than data cache
- ❑ Cache write:
 - How do we keep data in the cache and memory consistent?
- ❑ Two options (decision time again :-)
 - Write Back: write to cache only. Write the cache block to memory when that cache block is being replaced on a cache miss.
 - Use a “dirty” bit for each cache block
 - Multiple writes within a cache block require one write to memory
 - Greatly reduces the memory bandwidth requirement
 - Control can be complex
 - Write Through: write to cache and memory at the same time.
 - Simplifies data coherency
 - What!!! How can this be? Isn't memory too slow for this?

Write Buffer for Write Through



- ❑ A Write Buffer is placed between the Cache and Memory
 - Processor: writes data into the cache and the write buffer
 - Memory controller: write contents of the buffer to memory
- ❑ Write buffer is just a FIFO queue:
 - Typical number of entries: 4
 - Works fine if: Store frequency (w.r.t. time) $\ll 1 / \text{DRAM write cycle}$
- ❑ Memory system designer's nightmare:
 - Store frequency (w.r.t. time) $> 1 / \text{DRAM write cycle}$
 - Or bursty writes that exceed DRAM write speed
 - Write buffer saturation

Big Idea




- ❑ Create the illusion of a memory that is large, cheap, and fast - on average
- ❑ How to choose between associativity, block size, replacement & write policy?
- ❑ Design against a performance model
 - Minimize: *Average Memory Access Time*
= Hit Time X Hit Rate
+ Miss Penalty x Miss Rate
 - influenced by technology & program behavior
- ❑ How can we improve miss penalty?

Improving Miss Penalty



- ❑ When caches first became popular, Miss Penalty ~ 10 processor clock cycles
- ❑ But today, for a 2400 MHz Processor (0.4 ns per clock cycle) and 40 ns to go to DRAM
⇒ 100 processor clock cycles!



Solution: another cache between memory and the processor cache: Second Level (L2) Cache

Multilevel Cache Example

□ Given

- CPU base CPI = 1, clock rate = 4GHz
- Miss rate/instruction = 2%
- Main memory access time = 100ns

□ With just primary cache

- Miss penalty = $100\text{ns} / 0.25\text{ns} = 400$ cycles
- Effective CPI = $1 + 0.02 \times 400 = 9$

Multilevel Cache Example (cont.)

❑ Now add L-2 cache

- Access time = 5ns
- Global miss rate to main memory = 0.5%

❑ Primary miss with L-2 hit

- Penalty = $5\text{ns}/0.25\text{ns} = 20$ cycles

❑ Primary miss with L-2 miss

- Extra penalty = 400 cycles

❑ $\text{CPI} = 1 + 0.02 \times 20 + 0.005 \times 400 = 3.4$

❑ $\text{Performance ratio} = 9/3.4 = 2.6$

Multilevel Cache Considerations

❑ Primary cache

- Focus on minimal hit time

❑ L-2 cache

- Focus on low miss rate to avoid main memory access
- Hit time has less overall impact

❑ Results




- L-1 cache usually smaller than a single cache
- L-1 block size smaller than L-2 block size

Interactions with Advanced CPUs

- ❑ Out-of-order CPUs can execute instructions during cache miss
 - Pending store stays in load/store unit
 - Dependent instructions wait
 - Independent instructions continue
- ❑ Effect of miss depends on program data flow
 - Much harder to analyze
 - Use system simulation

Interactions with Software

- ❑ Misses depend on memory access patterns
 - Algorithm behavior
 - Compiler optimization for memory access
- ❑ Understanding of the memory hierarchy is critical to understanding the performance of programs



And in Conclusion...

- ❑ We've discussed memory caching in detail.
Caching in general shows up over and over in computer systems
 - File system cache
 - Web page cache
- ❑ Big idea: if something is expensive but we want to do it repeatedly, do it once and cache the result.

And in Conclusion...

- ❑ Cache design choices:
 - size of cache: speed v. capacity
 - direct-mapped v. associative
 - for N-way set assoc: choice of N
 - Write through v. write back
 - block replacement policy
 - 2nd level cache? 3rd level cache?
- ❑ Use performance model to pick between choices, depending on programs, technology, budget, ...