
C335

Computer Structures

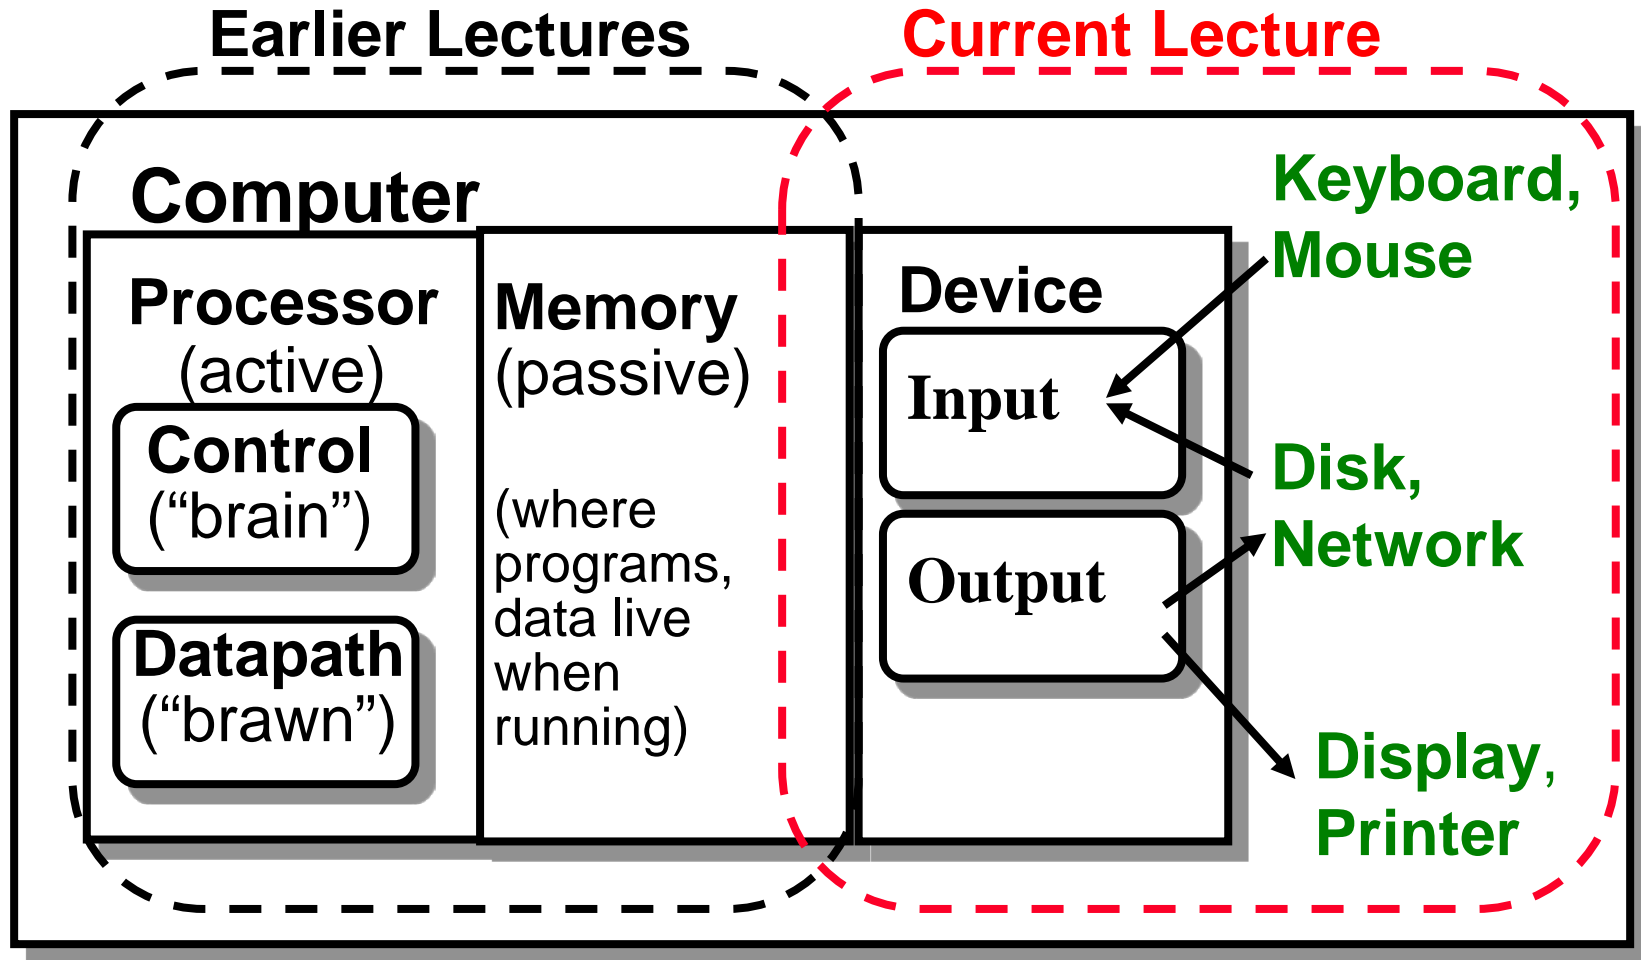
Input/Output

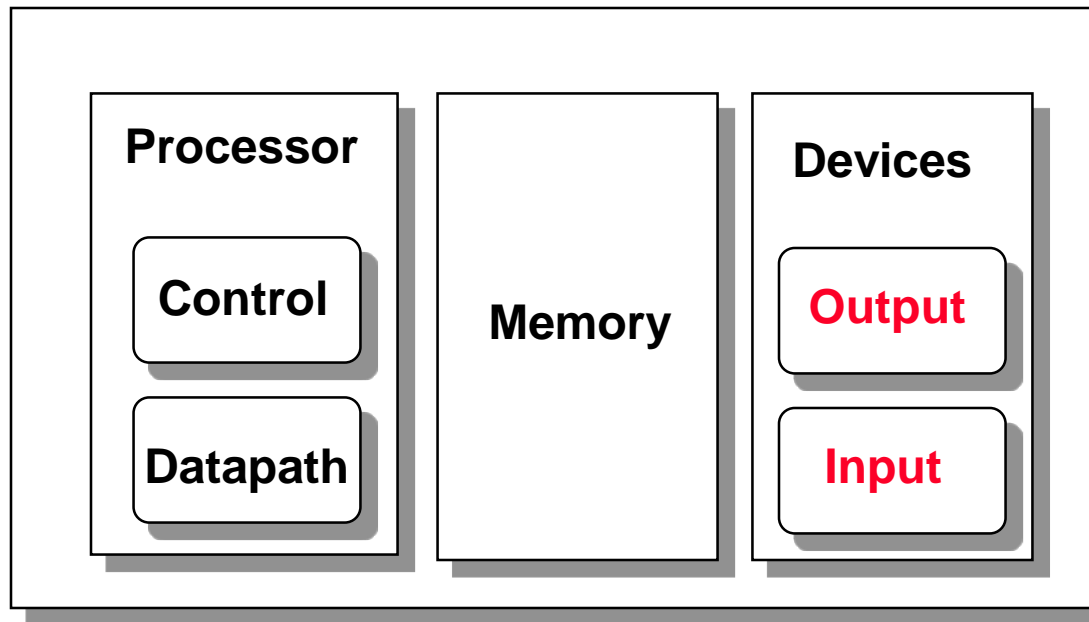

– Dealing with Exceptions and Interrupts

Dr. Liqiang Zhang

Department of Computer and Information Sciences

Recall : 5 components of any Computer





❑ Important metrics for an I/O system

- Dependability
- Expandability and diversity
- Performance
- Cost, size, weight

Input and Output Devices



❑ I/O devices are incredibly diverse with respect to

- Behavior – input, output or storage
- Partner – human or machine
- Data rate – the peak rate at which data can be transferred between the I/O device and the main memory or processor

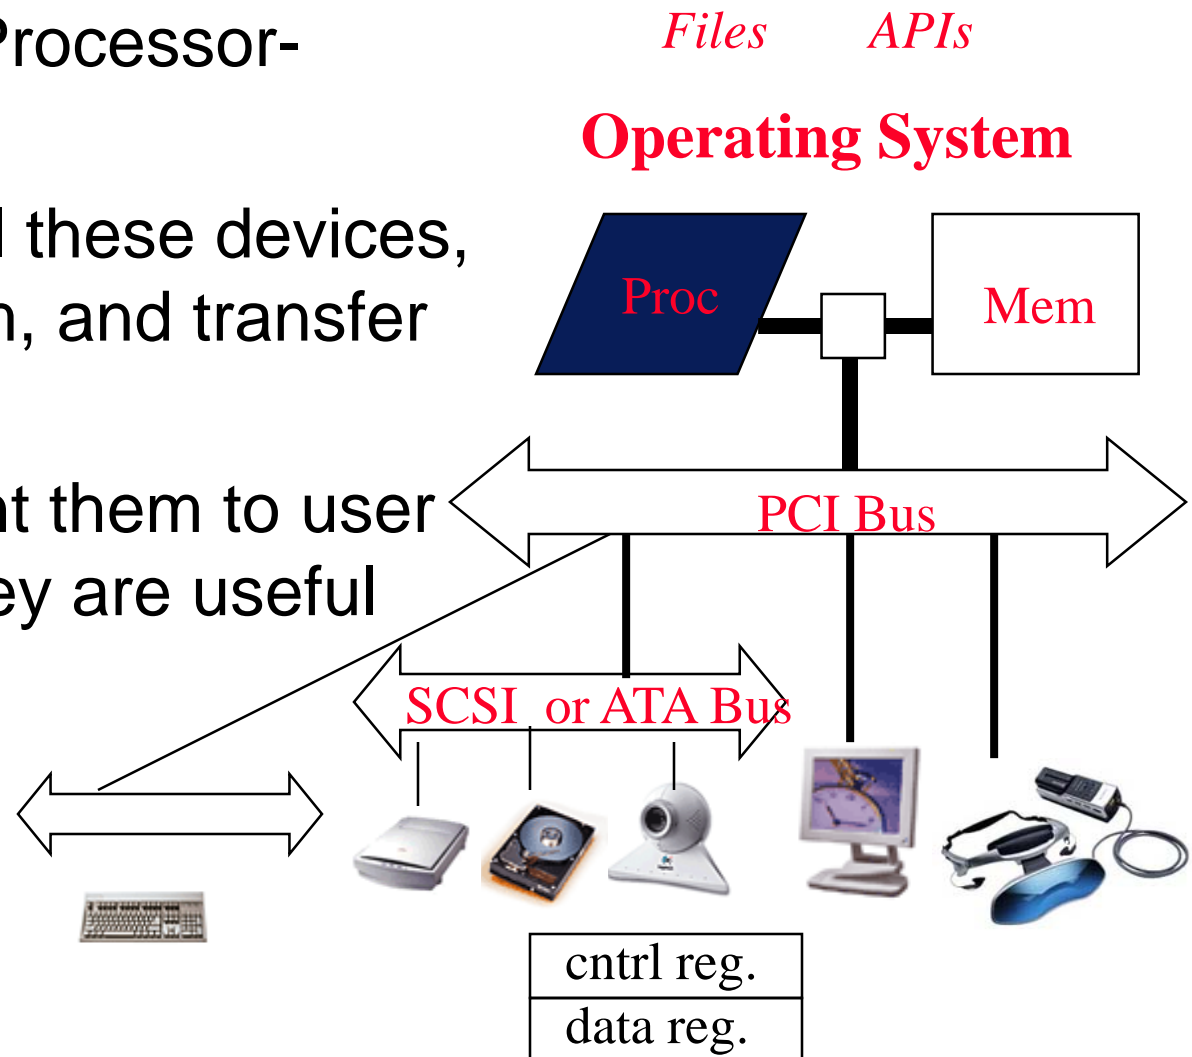
Device	Behavior	Partner	Data rate (Mb/s)
Keyboard	input	human	0.0001
Mouse	input	human	0.0038
Laser printer	output	human	3.2000
Graphics display	output	human	800.0000-8000.0000
Network/LAN	input or output	machine	100.0000-10000.0000
Magnetic disk	storage	machine	240.0000-2560.0000

8 orders of magnitude
range

What do we need to make I/O work?



- ❑ A way to connect many types of devices to the Processor-Memory
- ❑ A way to control these devices, respond to them, and transfer data
- ❑ A way to present them to user programs so they are useful





❑ How the processor directs the I/O devices

- Special I/O instructions (Standard I/O)
 - Must specify both the device and the command
 - Additional pin (*M/I/O*) on bus indicates whether a memory or peripheral access
- Memory-mapped I/O
 - Portions of the high-order memory address space are assigned to each I/O device
 - Read and writes to those memory addresses are interpreted as commands to the I/O devices

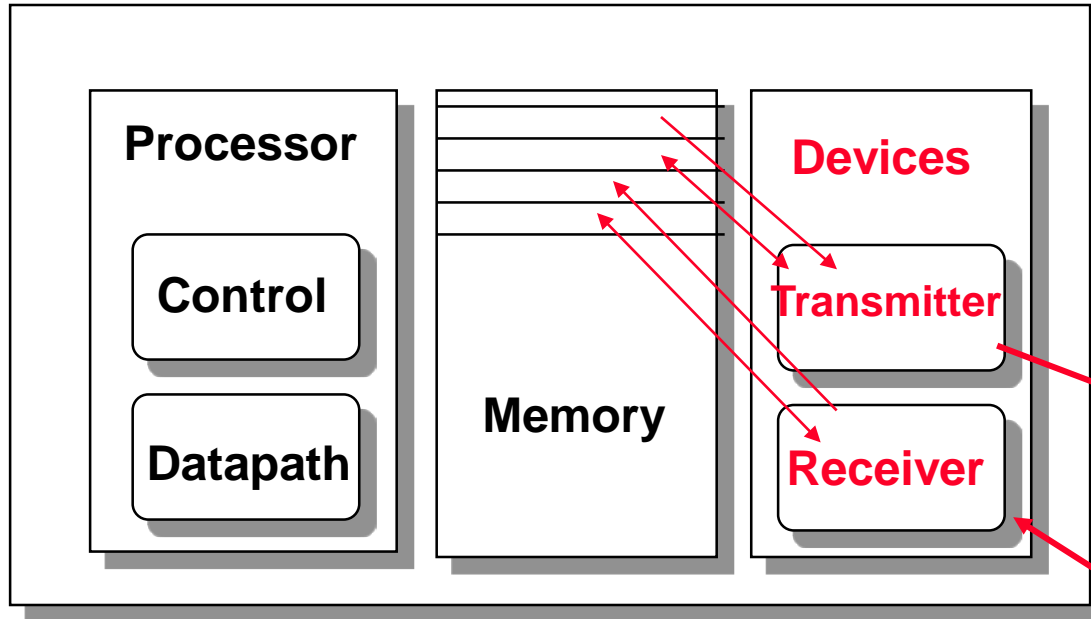
❑ How the I/O device communicates with the processor

- Polling – the processor periodically checks the status of an I/O device to determine its need for service
 - Processor is totally in control – but does **all** the work
 - Can waste a lot of processor time due to speed differences
- Interrupt-driven – the I/O device issues an interrupts to the processor to indicate that it needs attention

I/O in SPIM




- ❑ SPIM supports one memory-mapped I/O device – a terminal with two independent units
 - Receiver reads characters from the keyboard
 - Transmitter writes characters to the display



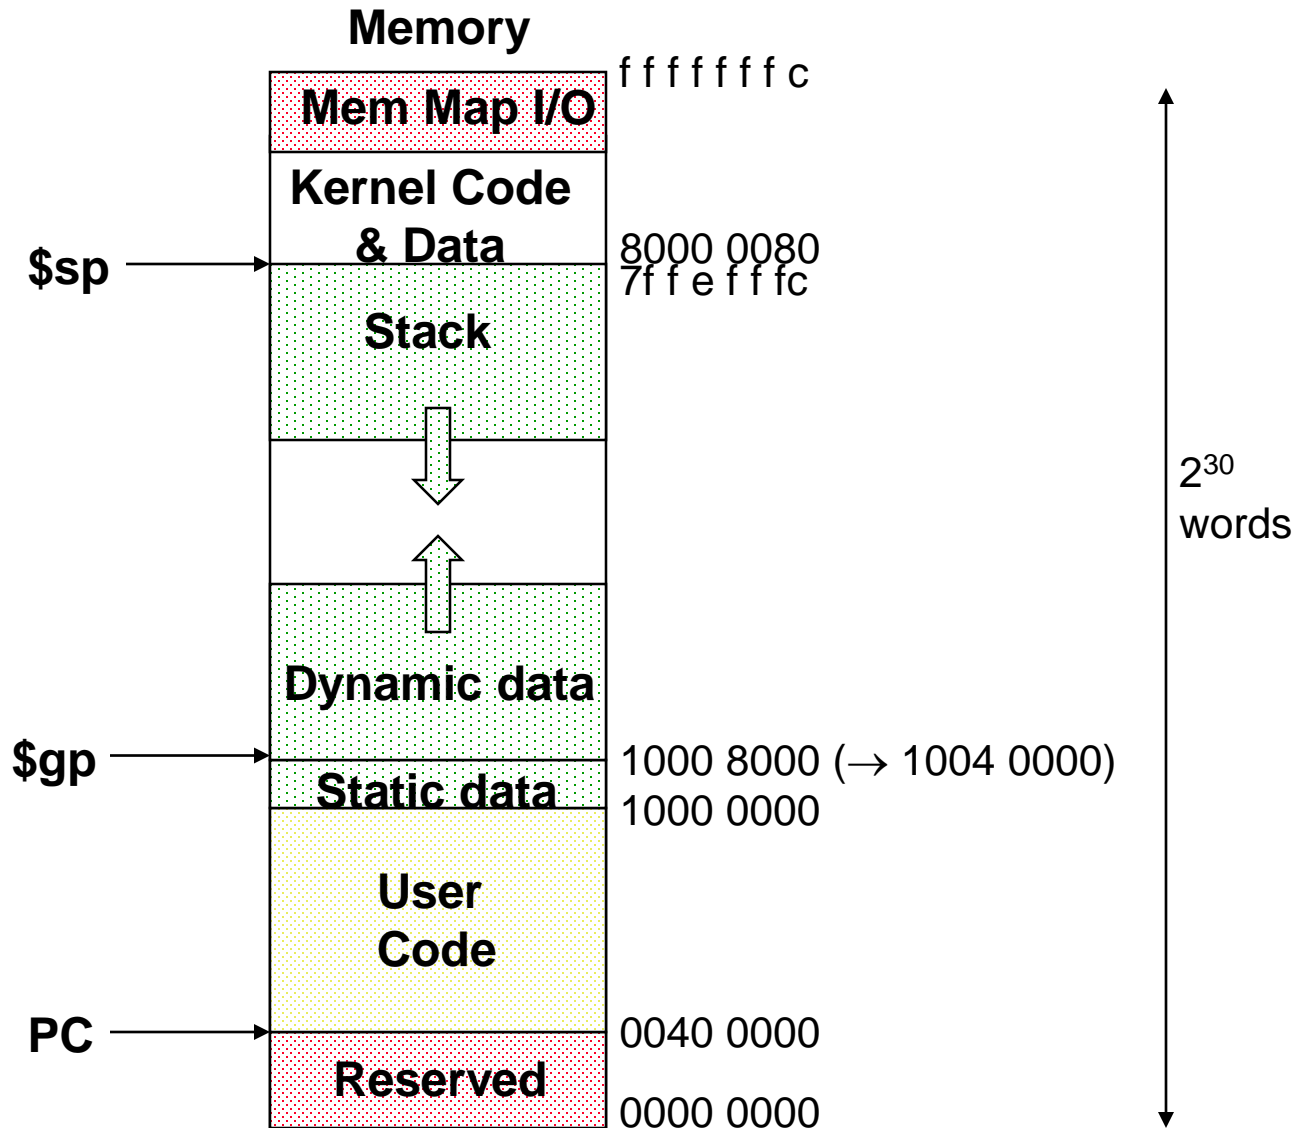
Memory Mapped I/O



- ❑ Certain addresses are not regular memory
- ❑ Instead, they correspond to registers in I/O devices

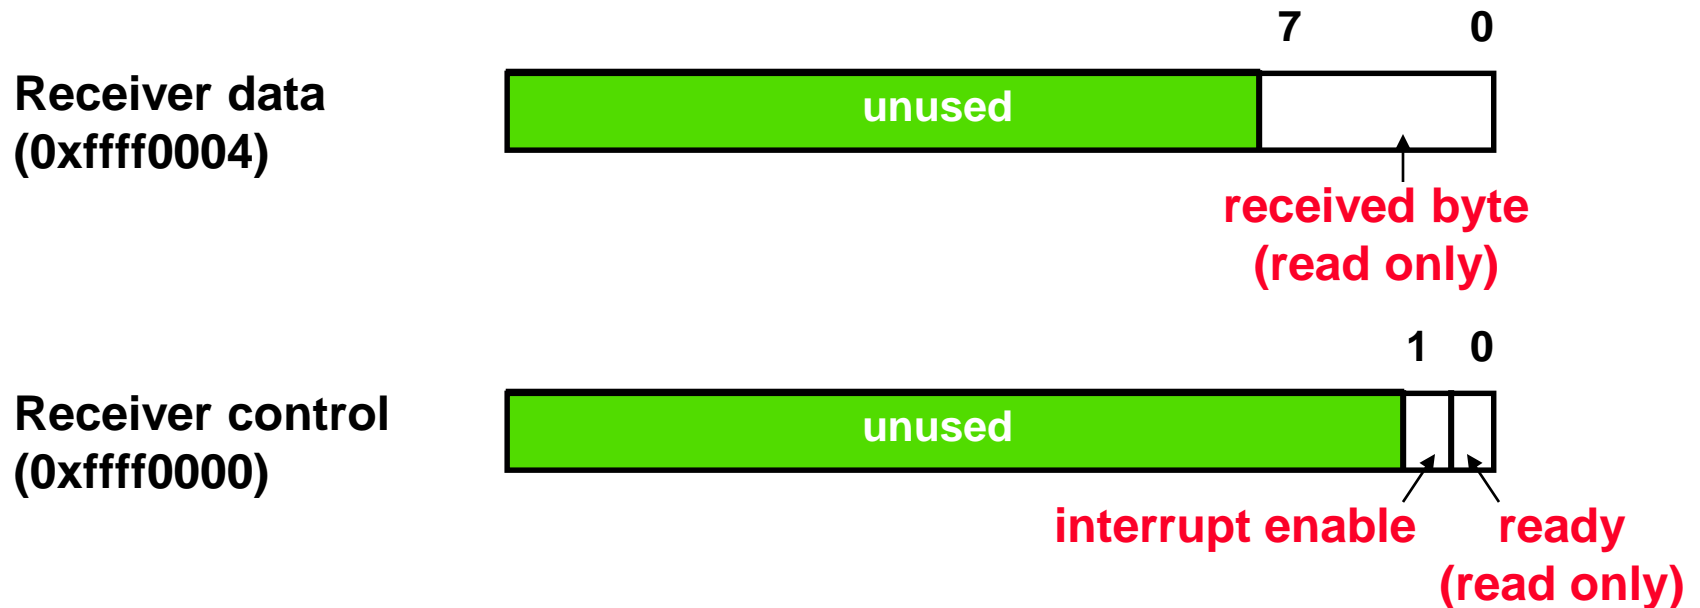


Review: MIPS (spim) Memory Allocation




Terminal Receiver (Input) Control with SPIM

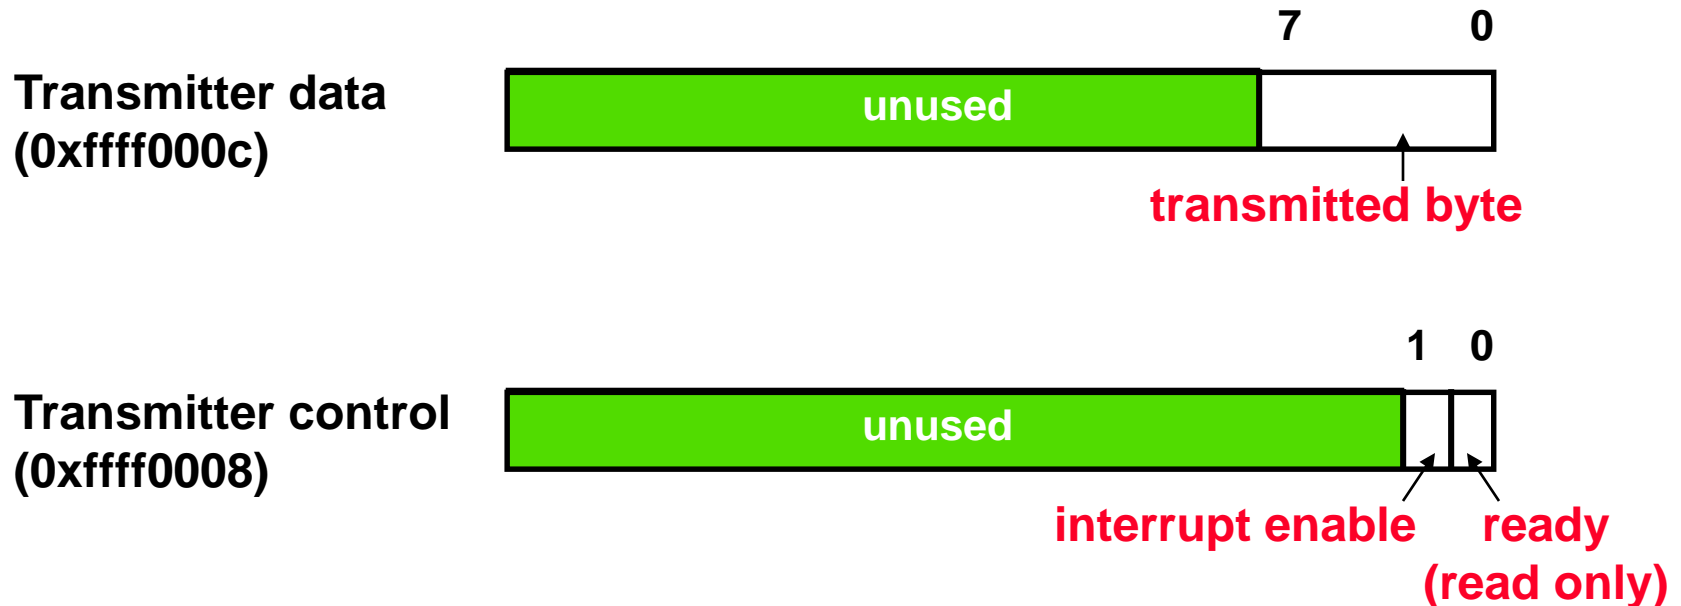
- ❑ Terminal input is controlled via two memory-mapped device registers (i.e., each register appears as a special memory location)



Terminal Output Control with SPIM



- ❑ Terminal output is controlled via two memory-mapped device registers (i.e., each register appears as a special memory location)



Processor-I/O Speed Mismatch



- ❑ 1GHz microprocessor can execute 1 billion load or store instructions per second, or 32,000Mb/s data rate
 - I/O devices data rates range from 0.0001 Mb/s to 10,000 Mb/s
- ❑ Input: device may not be ready to send data as fast as the processor loads it
 - Also, might be waiting for human to act
- ❑ Output: device may not be ready to accept data as fast as processor stores it
- ❑ What to do?

Processor Checks Status before Acting



- ❑ Path to device generally has 2 registers:
 - Control Register, says it's OK to read/write (I/O ready) [think of a flagman on a road]
 - Data Register, contains data
- ❑ Processor reads from Control Register in loop, waiting for device to set Ready bit in Control reg (0 \Rightarrow 1) to say its OK
- ❑ Processor then loads from (input) or writes to (output) data register
 - Load from or Store into Data Register resets Ready bit (1 \Rightarrow 0) of Control Register



❑ Control register rightmost bit (0): Ready

- Receiver: Ready==1 means character in Data Register not yet been read;
1 \Rightarrow 0 when data is read from Data Reg
- Transmitter: Ready==1 means transmitter is ready to accept a new character;
0 \Rightarrow Transmitter still busy writing last char

❑ Data register rightmost byte has data

- Receiver: last char from keyboard; rest = 0
- Transmitter: when write rightmost byte, writes char to display

I/O Example



- ❑ Input: Read from keyboard into \$v0

```
Waitloop:      lui    $t0, 0xffff #ffff0000
               lw     $t1, 0($t0) #control
               andi   $t1,$t1,0x1
               beq    $t1,$zero, Waitloop
               lw     $v0, 4($t0) #data
```

- ❑ Output: Write to display from \$a0

```
Waitloop:      lui    $t0, 0xffff #ffff0000
               lw     $t1, 8($t0) #control
               andi   $t1,$t1,0x1
               beq    $t1,$zero, Waitloop
               sw     $a0, 12($t0) #data
```

- ❑ “Ready” bit is from processor’s point of view!
- ❑ Processor waiting for I/O called “Polling”

Cost of Polling?



- ❑ Assume for a processor with a 1GHz clock it takes 400 clock cycles for a polling operation (call polling routine, accessing the device, and returning). Determine % of processor time for polling
 - Mouse: polled 30 times/sec so as not to miss user movement
 - Floppy disk: transfers data in 2-Byte units and has a data rate of 50 KB/second.
No data transfer can be missed.
 - Hard disk: transfers data in 16-Byte chunks and can transfer at 16 MB/second. Again, no transfer can be missed.

% Processor time to poll



Mouse Polling [clocks/sec]

$$= 30 \text{ [polls/s]} * 400 \text{ [clocks/poll]} = 12\text{K [clocks/s]}$$

❑ % Processor for polling:

$$12 * 10^3 \text{ [clocks/s]} / 1 * 10^9 \text{ [clocks/s]} = 0.0012\%$$

❑ Polling mouse little impact on processor

Frequency of Polling Floppy

$$= 50 \text{ [KB/s]} / 2 \text{ [B/poll]} = 25\text{K [polls/s]}$$

❑ Floppy Polling, Clocks/sec

$$= 25\text{K [polls/s]} * 400 \text{ [clocks/poll]} = 10\text{M [clocks/s]}$$

❑ % Processor for polling:

$$10 * 10^6 \text{ [clocks/s]} / 1 * 10^9 \text{ [clocks/s]} = 1\%$$

❑ OK if not too many I/O devices

% Processor time to poll hard disk



Frequency of Polling Disk

$$= 16 \text{ [MB/s]} / 16 \text{ [B/poll]} = 1\text{M [polls/s]}$$

❑ Disk Polling, Clocks/sec

$$= 1\text{M [polls/s]} * 400 \text{ [clocks/poll]}$$

$$= 400\text{M [clocks/s]}$$

❑ % Processor for polling:

$$400 * 10^6 \text{ [clocks/s]} / 1 * 10^9 \text{ [clocks/s]} = 40\%$$

⇒ Unacceptable

What is the alternative to polling?

- ❑ Wasteful to have processor spend most of its time “spin-waiting” for I/O to be ready
- ❑ Would like an unplanned procedure call that would be invoked only when I/O device is ready
- ❑ Solution: use **exception mechanism** to help I/O. **Interrupt** program when I/O ready, return when done with data transfer



- ❑ An I/O interrupt is like overflow exceptions except:
 - An I/O interrupt is “asynchronous”
 - More information needs to be conveyed
- ❑ An I/O interrupt is asynchronous with respect to instruction execution:
 - I/O interrupt is not associated with any instruction, but it can happen in the middle of any given instruction
 - I/O interrupt does not prevent any instruction from completion

I/O Interrupts



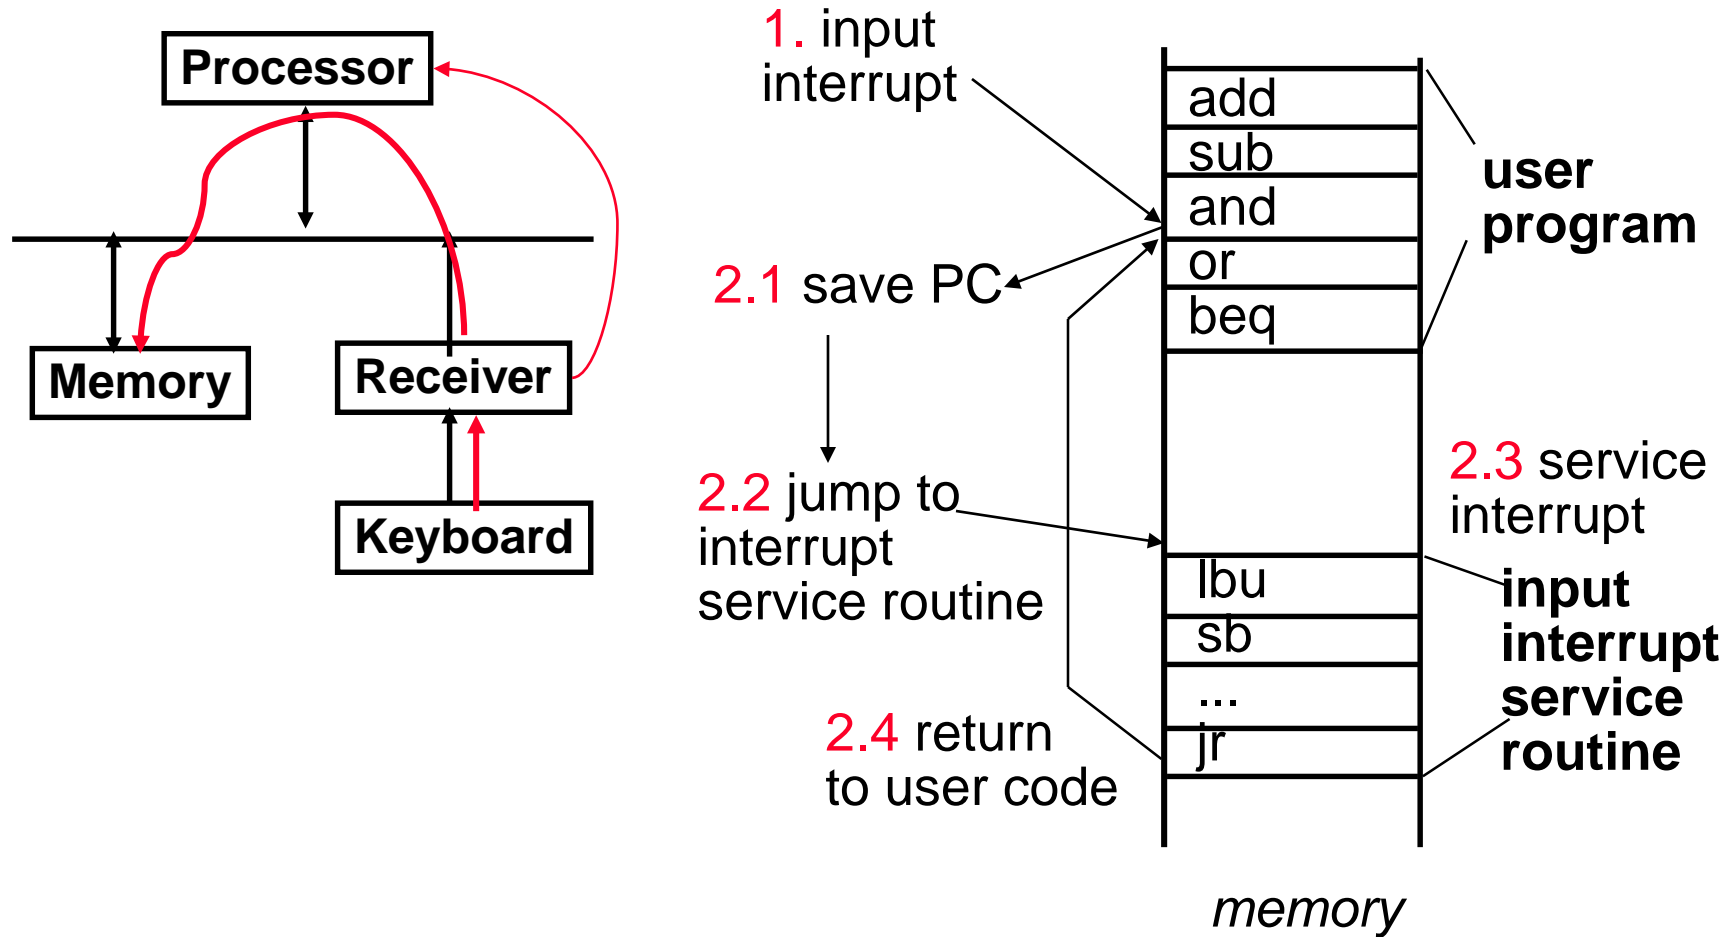
- ❑ An I/O interrupt is used to signal an I/O request for service
 - Can have different urgencies (so may need to be prioritized)
 - Need to somehow convey the identity of the device generating the interrupt
- ❑ Advantage
 - User program progress is only halted during actual transfer of I/O data to/from user memory space
- ❑ Disadvantage - special hardware is needed to
 - Cause an interrupt (I/O device)
 - Detect an interrupt and save the proper information to resume after servicing the interrupt (processor)

Benefit of Interrupt-Driven I/O



- ❑ Find the % of processor consumed if the hard disk is only active 5% of the time. Assuming 500 clock cycle overhead for each transfer, including interrupt:
 - Disk Interrupts/s = $16 \text{ [MB/s]} / 16 \text{ [B/interrupt]}$
= 1M [interrupts/s]
 - Disk Interrupts [clocks/s]
= $1\text{M [interrupts/s]} * 500 \text{ [clocks/interrupt]}$
= $500,000,000 \text{ [clocks/s]}$
 - % Processor for during transfer:
 $500 * 10^6 \text{ [clocks/s]} / 1 * 10^9 \text{ [clocks/s]} = 50\%$
- ❑ Disk active 5% $\Rightarrow 5\% * 50\% \Rightarrow 2.5\%$ busy

Interrupt Driven Input



Interrupt Driven Input in SPIM

1. the Receiver indicates with an **interrupt** that it has input a new character from the keyboard into the Receiver data register

Receiver data
(0xffff0004)



- writing to the Receiver data register sets the Receiver control register ready bit to 1

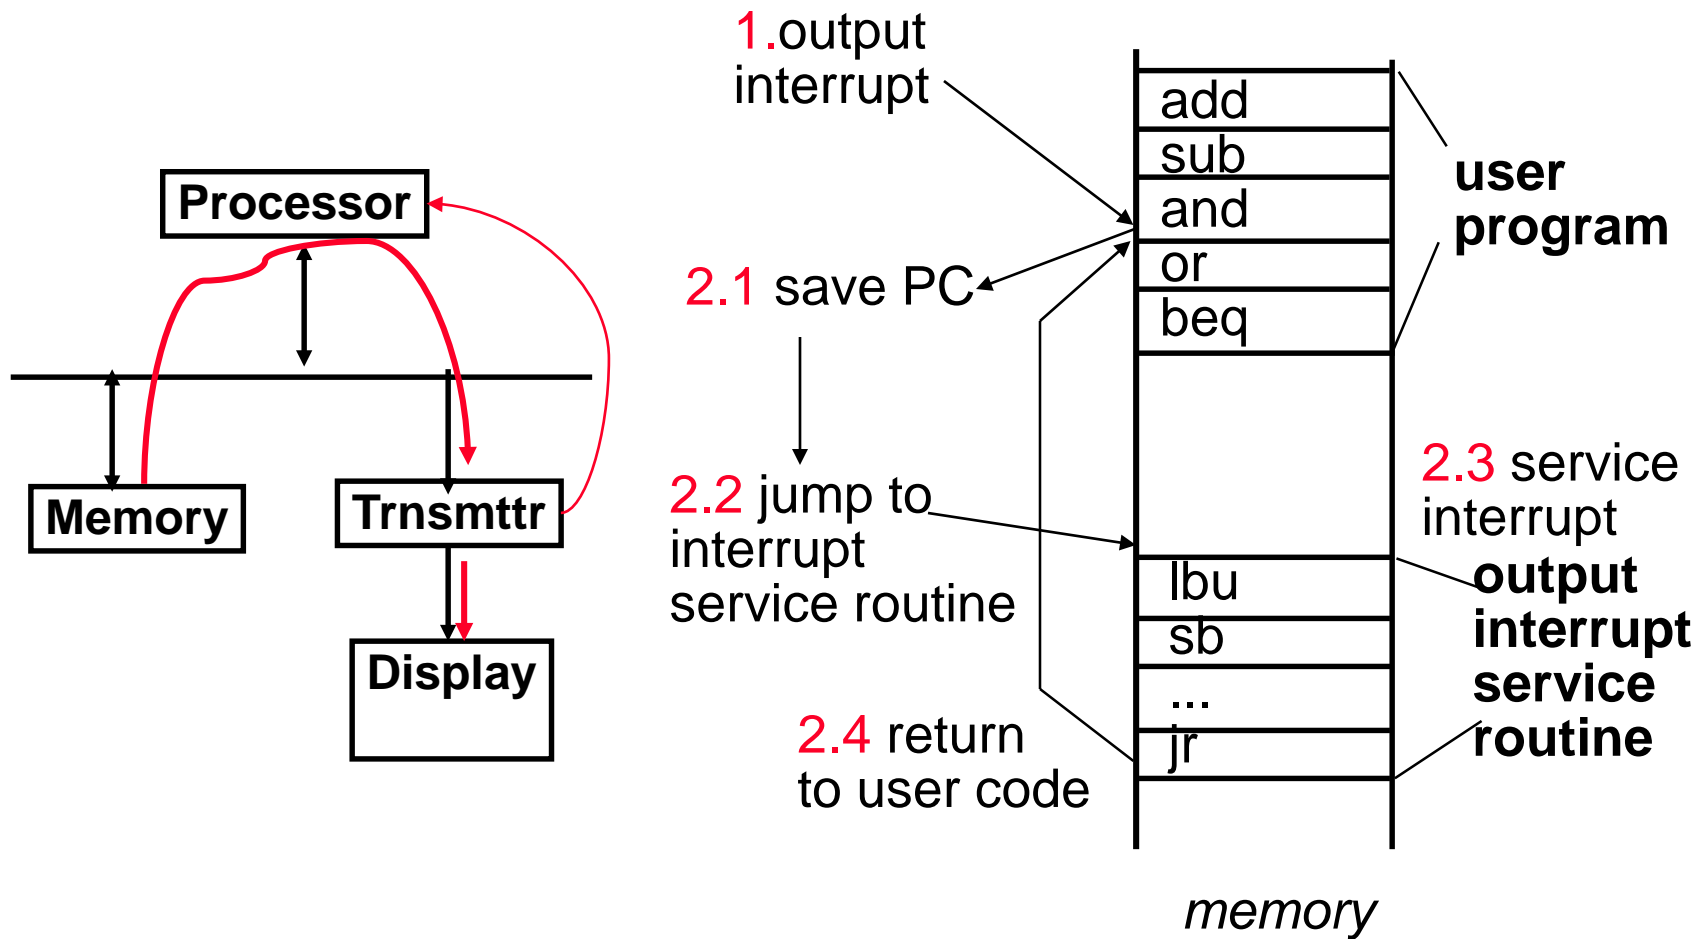
Receiver control
(0xffff0000)



2. the user process responds to the **interrupt** by transferring control to an interrupt service routine that copies the input character into the user memory space

- reading the Receiver data register resets the Receiver control register ready bit to 0

Interrupt Driven Output



Interrupt Driven Output in SPIM



1. the transmitter indicates with an **interrupt** that it has successfully output the character in the Transmitter data register in memory to the display **transmitted byte**

Transmitter data
(0xffff000c)



- reading from the Transmitter data register sets the Transmitter control register ready bit to 1

Transmitter control
(0xffff0008)

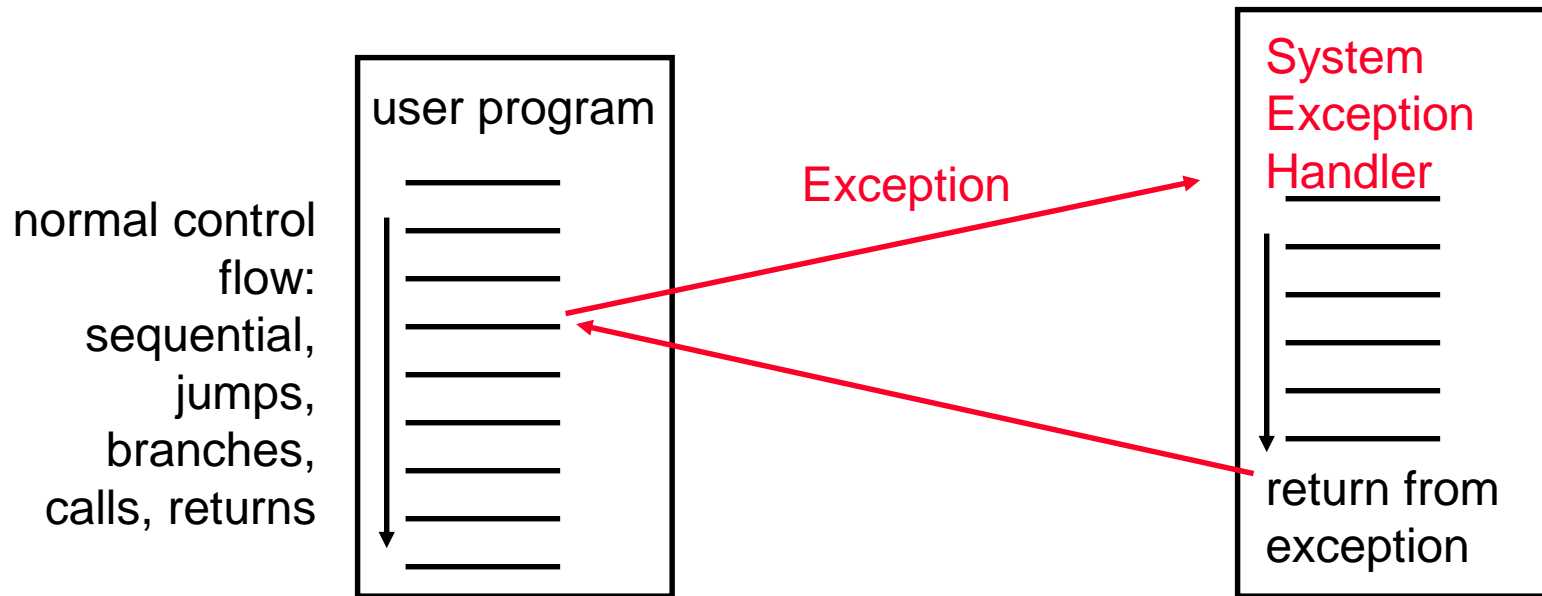



interrupt enable **ready**

2. the user process responds to the **interrupt** by transferring control to an interrupt service routine that writes the next character to output from the user memory space into the Transmitter data register

- writing to the Transmitter data register resets the Transmitter control register ready bit to 0

Exceptions in General



□ Exception = **unprogrammed** control transfer

- system takes action to handle the exception
 - must record the address of the offending or next to execute instruction and save (and restore) user state
- returns control to user after handling the exception

Two Types of Exceptions



❑ Interrupts

- caused by **external events** (i.e., request from I/O device)
- **asynchronous** to program execution
- may be handled **between** instructions
- simply suspend and resume user program


❑ Traps (Exception)

- caused by **internal events**
 - exceptional conditions (e.g., arithmetic overflow, undefined instr.)
 - errors (e.g., hardware malfunction, memory parity error)
 - faults (e.g., non-resident page – page fault)
- **synchronous** to program execution
- condition must be remedied by the trap handler
- instruction may be retried and program continued or program may be aborted

Additions to MIPS ISA for I/O Interrupts

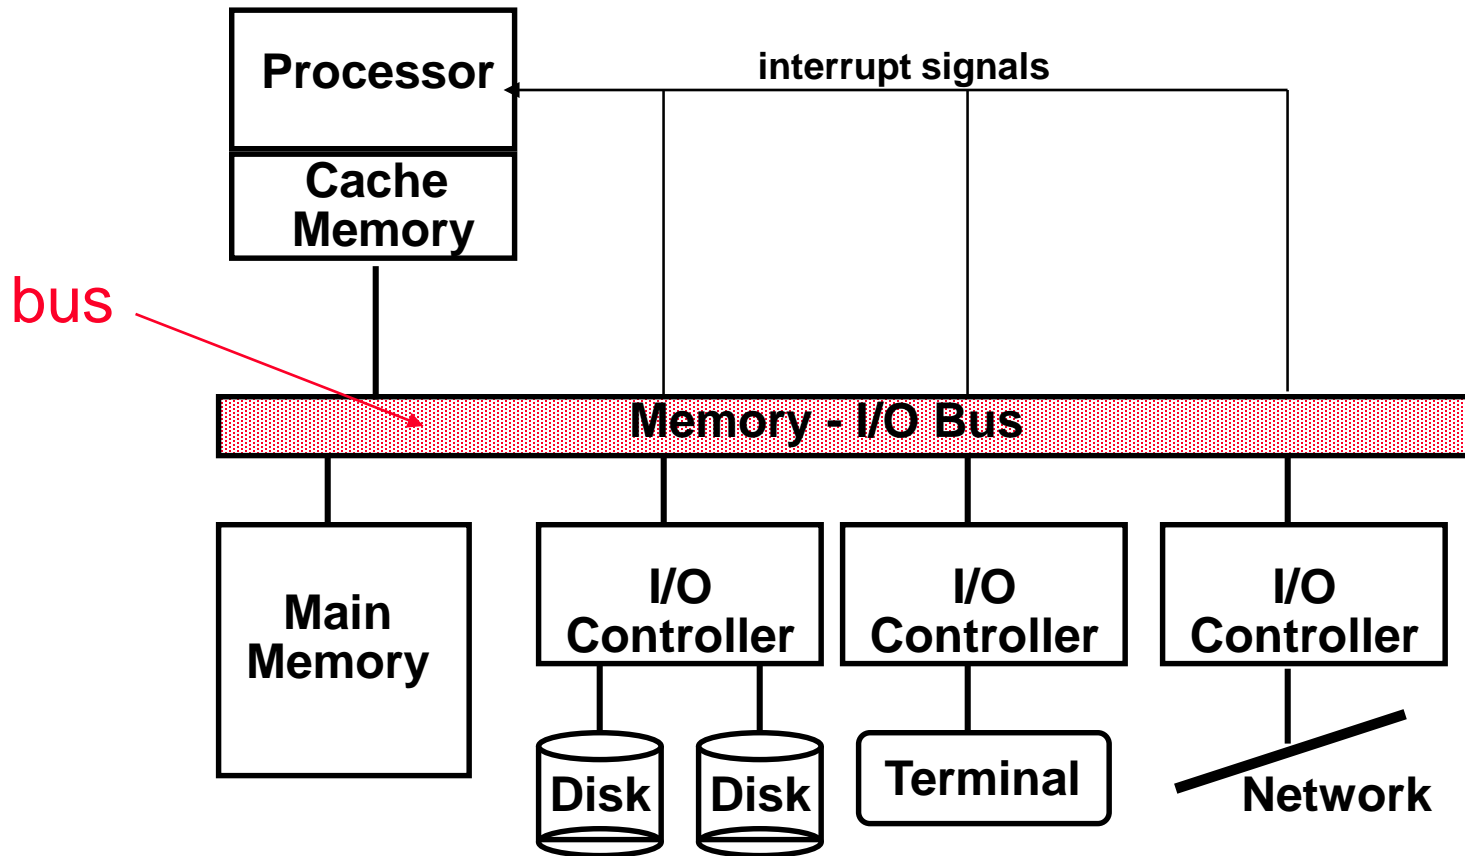
- ❑ Control signals to write EPC, Cause, and Status
- ❑ Hardware to record the type of interrupt in Cause
- ❑ Modify the Datapath and Control, so that
 - the address of **interrupt handler** ($8000\ 0080_{\text{hex}}$) can be loaded into the PC, so must increase the size of PC multiplex
 - and save the address of the current/next instr in EPC

Example I/O Interrupt Service Routine



```
        .ktext 0x80000080
intrpt: sw      $t0, save0           #save working reg's
        . . .
        mfc0    $k0, $13           #move Cause into $k0
        mfc0    $k1, $14           #move EPC into $k1
        sgt     $t0, $k0, 0x00     #is it an interrupt?
        bgtz    $t0, excpt
        li      $t0, 0xffffffff
        lw      $t1, 0($t0)        #test Receiver control
        beq     $t1, $0, tstout    # for input ready
inpt:    lb      $v0, 4($t0)        #input char. into $v0
tstout:  lw      $t1, 8($t0)        #test Trans. control
        beq     $t1, $0, excpt    # for output ready
        sb      $v1, c($t0)       #output char. from $v1
        j       done              #
excpt:   . . .                     #code to handle excpts
        addiu   $k1, $k1, 4        #do not re-execute
        . . .                     # faulting instr
done:    lw      $t0, save0        #restore working reg's
        . . .
        eret                       #return to user code (EPC)
        .kdata
save0:   .word   0
        . . .
```

I/O System Interconnect Issues



- ❑ Usually have more than one I/O device in the system connected to the processor via a bus
 - each I/O device is controlled by an I/O Controller

Buses



- ❑ A **bus** is a shared communication link (a single set of wires used to connect multiple subsystems) that needs to support a range of devices with widely varying latencies and data transfer rates
 - Advantages
 - Versatile – new devices can be added easily and can be moved between computer systems that use the same bus standard
 - Low cost – a single set of wires is shared in multiple ways
 - Disadvantages
 - Creates a communication bottleneck – bus **bandwidth** limits the maximum I/O **throughput**
- ❑ The maximum bus speed is largely limited by
 - The **length** of the bus
 - The **number** of devices on the bus

I/O Performance Measures



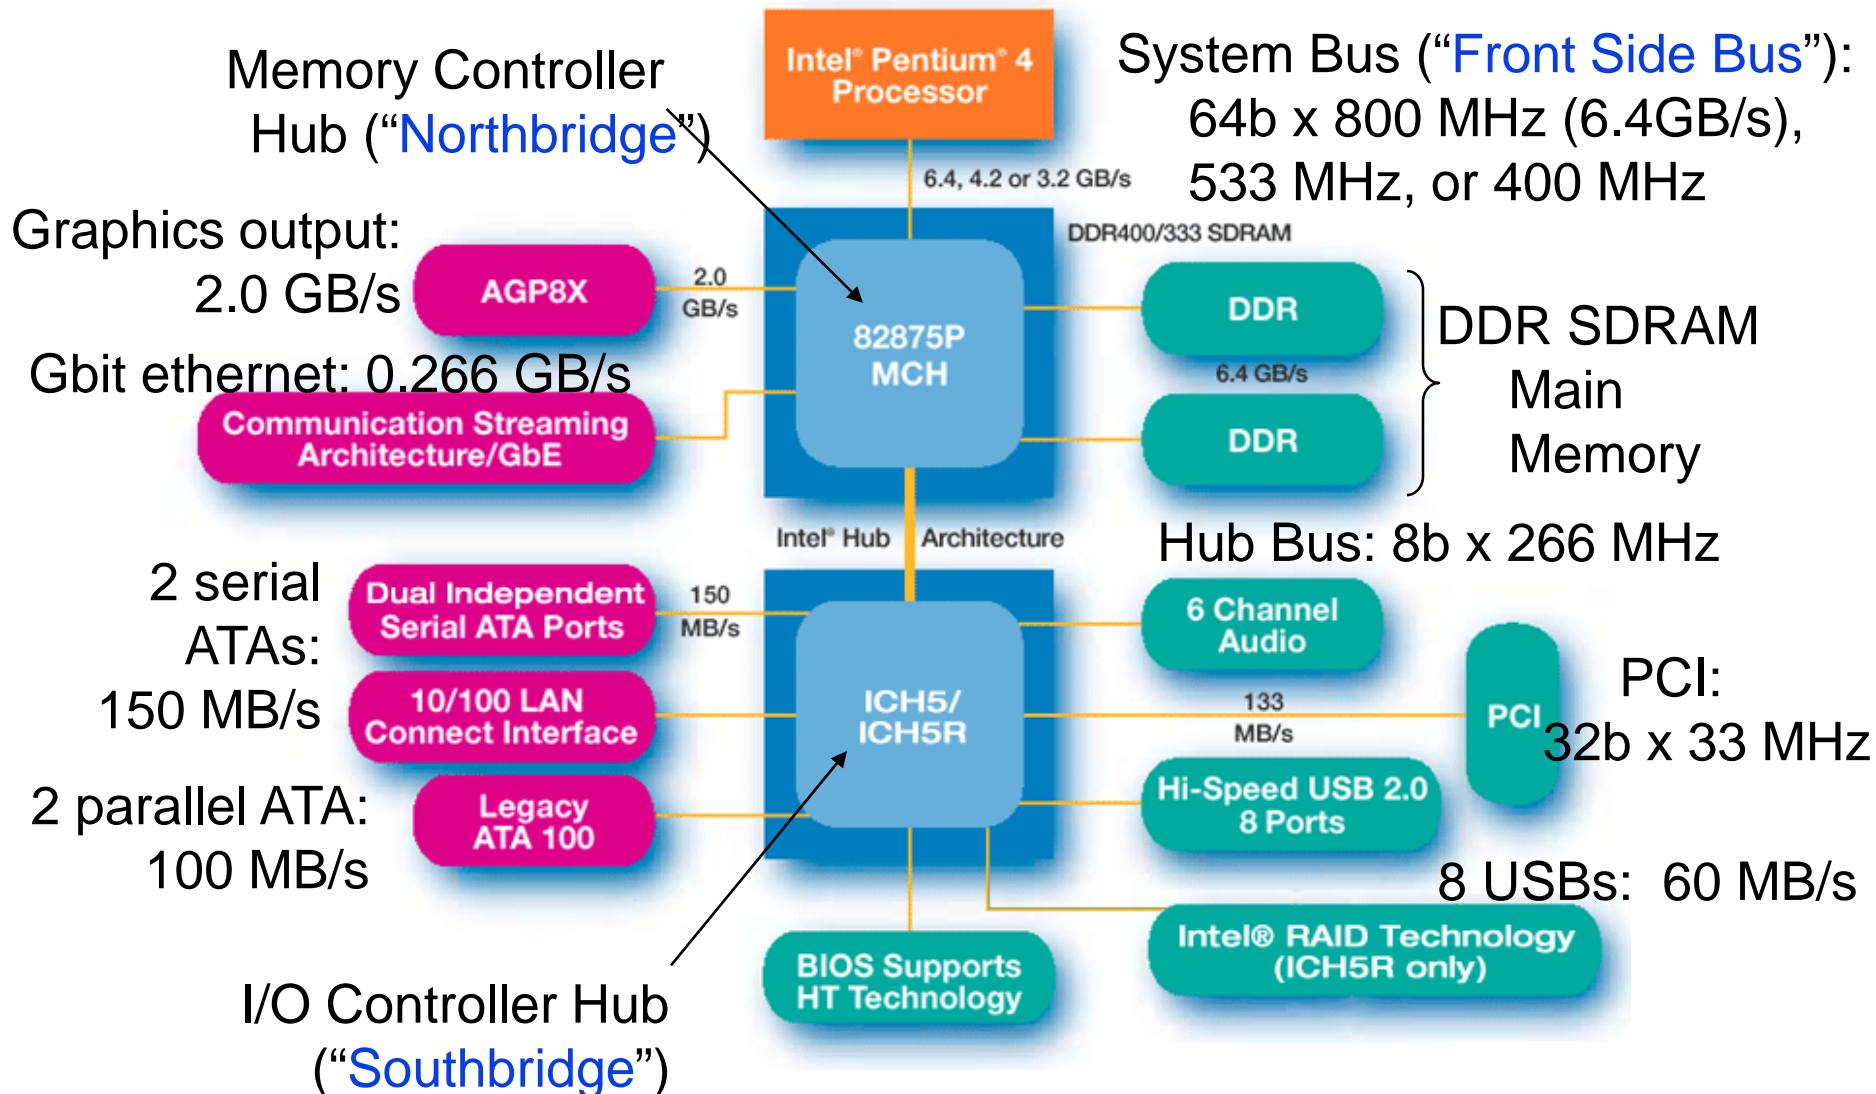
- ❑ **I/O bandwidth** (throughput) – amount of information that can be input (output) and communicated across an interconnect (e.g., a bus) to the processor/memory (I/O device) per unit time
 1. How much data can we move through the system in a certain time?
 2. How many I/O operations can we do per unit time?
- ❑ **I/O response time** (latency) – the total elapsed time to accomplish an input or output operation
 - An especially important performance metric in real-time systems
- ❑ Many applications require both high throughput and short response times

Types of Buses

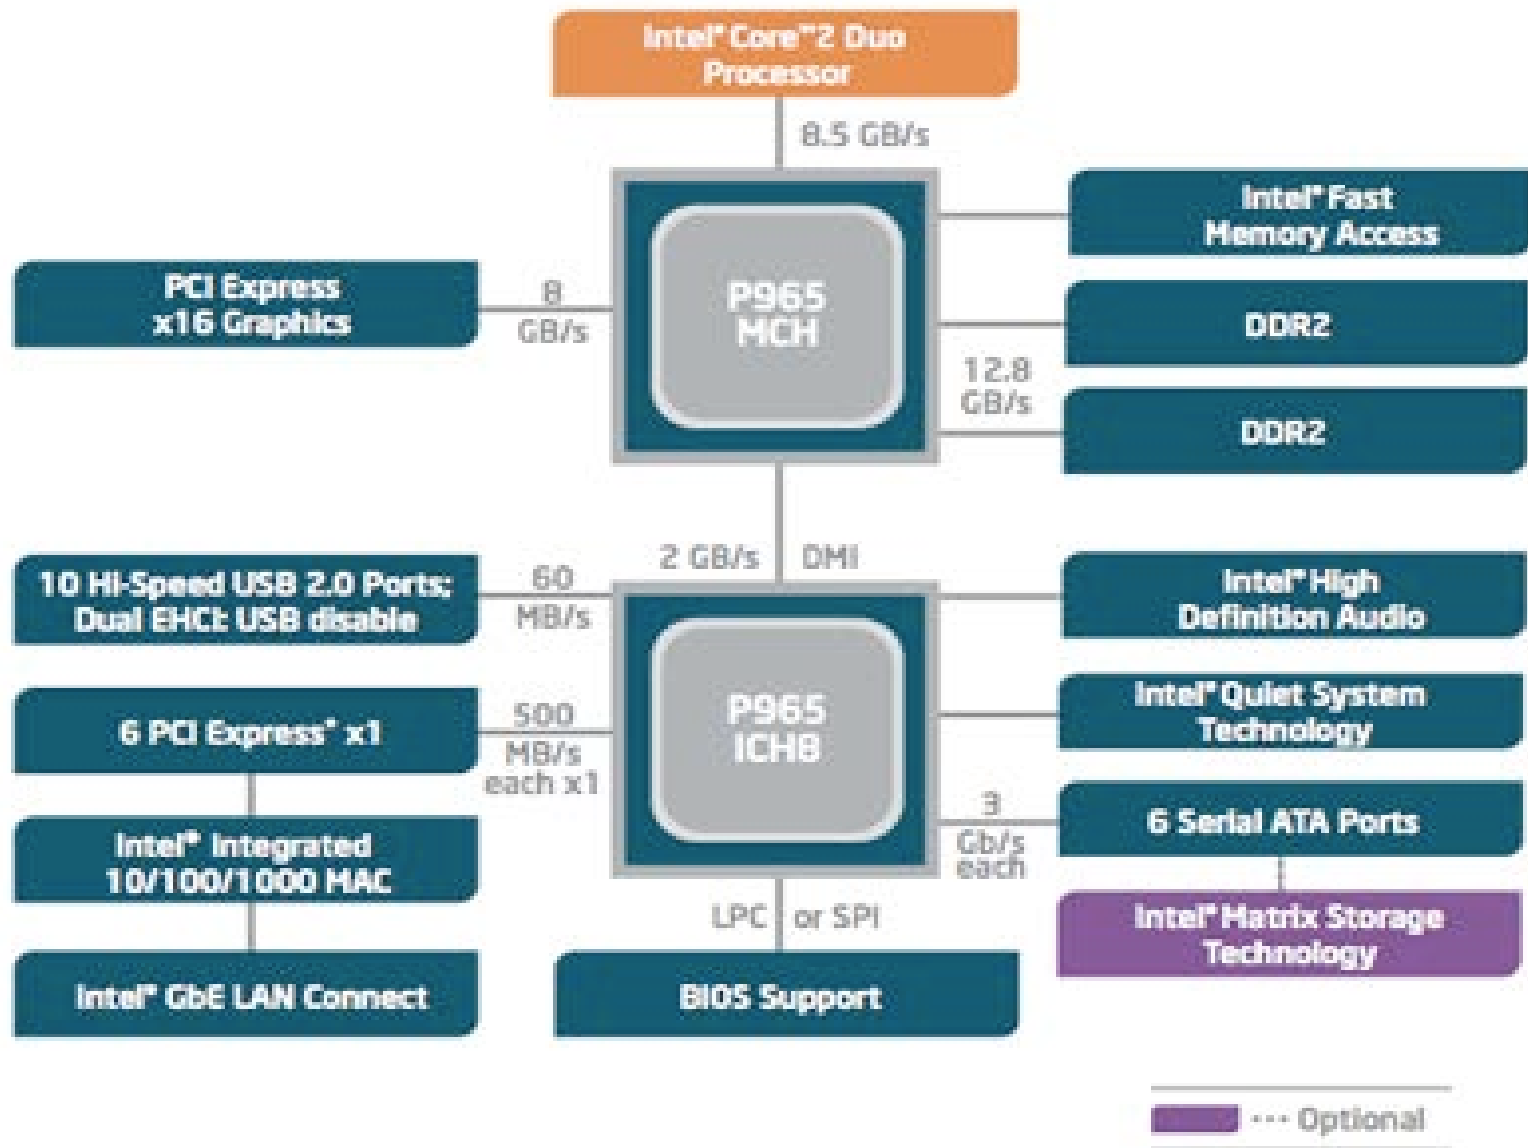


- ❑ Processor-memory bus (proprietary)
 - Short and high speed
 - Matched to the memory system to maximize the memory-processor bandwidth
 - Optimized for cache block transfers
- ❑ I/O bus (industry standard, e.g., SCSI, USB, Firewire)
 - Usually is lengthy and slower
 - Needs to accommodate a wide range of I/O devices
 - Connects to the processor-memory bus or backplane bus
- ❑ Backplane bus (industry standard, e.g., ISA, PCI, PCI Express)
 - The backplane is an interconnection structure within the chassis
 - Used as an intermediary bus connecting I/O busses to the processor-memory bus

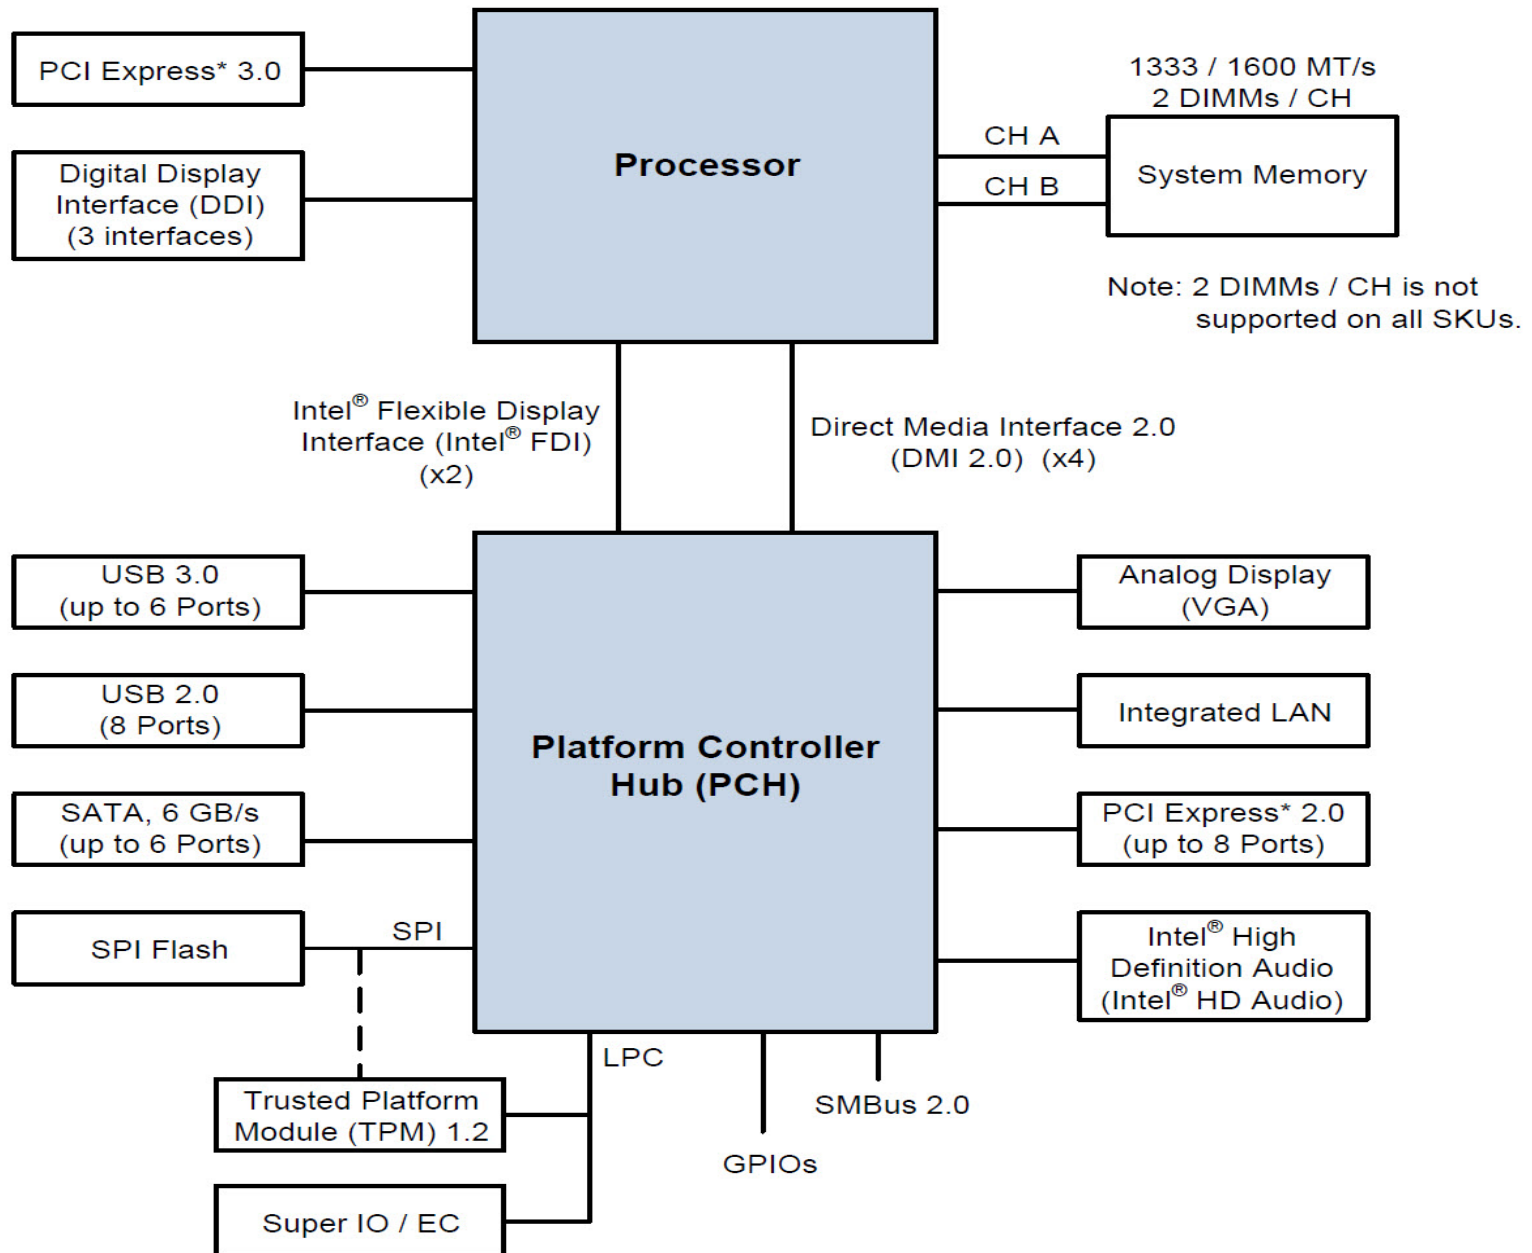
Example: The Pentium 4's Buses



Example: The Pentium Core 2 Duo Buses



Example: The Intel 4th Gen Core i7



Bus Bandwidth Determinates

- ❑ The bandwidth of a bus is determined by
 - Whether its is synchronous or asynchronous and the timing characteristics of the protocol used
 - The bus width (i.e., number of data lines)
 - Whether the bus supports block transfers or only word at a time transfers

	Firewire	USB 2.0
Type	I/O	I/O
Data lines	4	2
Architecture	Peer-to-peer	Master-slave
Max # devices	63	127
Max length	4.5 meters	5 meters
Peak bandwidth	400 Mbps , 800 Mbps 1600 Mbps, 3200 Mbps	1.5 Mb/s (low), 12 Mb/s (full), 480 Mb/s (high), 5Gbps (super)

Buses in Transition

- ❑ Companies are transitioning from synchronous, parallel, *wide* buses to asynchronous *narrow* buses
 - Reflection on wires and clock skew makes it difficult to use 16 to 64 parallel wires running at a high clock rate (e.g., ~400 MHz) so companies are transitioning to buses with a few one-way wires running at a very high “clock” rate (~2 GHz)

	PCI	PCI Express x1	ATA	Serial ATA
Total # wires	120	36	80	7
Clock (MHz)	33 – 133	16000	50	16000
Peak BW (MB/s)	128 – 1064	1969 (v4.0)	100	1969 (rev. 3.2)

Buses in Transition

